

Review Session

Memory Hierarchy, Systems Fundamentals, Virtual
Memory

How to prepare for the exam?

- Read through all the lecture slides
- Review all the weekly quiz
- Use office hours if you have any doubt about any concept

Memory Hierarchy

Things you should know from this chapter:

- Cache Miss vs Hit and Locality
- Tag bits, set bits, and offset bits in memory address
- Cache Placement Policy: Direct Mapped, Fully Associative, and E-way Set Associative Cache
- Replacement Algorithm: LRU and LFU

Cache Hit / Miss and Locality

- Cache Miss vs Cache Hit
- Temporal Locality vs Spatial Locality
 - A program with good locality can significantly reduce miss rate and their penalty
- Block is the basic unit when transfer between memory

Cache Hit / Miss and Locality

- Cache Miss vs Cache Hit
- Temporal Locality vs Spatial Locality
 - A program with good locality can significantly reduce miss rate and their penalty
- Block is the basic unit when transfer between memory

What is the minimal number of block loads from grid if the block size is 8 bytes?

```
int total = 0;
for (i = 0; i < 3; ++i) {
    for (j = 0; j < 3; ++j) {
        total = grid[i][j];
    }
}
```

Cache Hit / Miss and Locality

- Cache Miss vs Cache Hit
- Temporal Locality vs Spatial Locality
 - A program with good locality can significantly reduce miss rate and their penalty
- Block is the basic unit when transfer between memory

What is the minimal number of block loads from grid if the block size is 8 bytes?

Answer: 5

```
int total = 0;
for (i = 0; i < 3; ++i) {
    for (j = 0; j < 3; ++j) {
        total = grid[i][j];
    }
}
```

grid	(0, 0)	(0, 1)	(0, 2)	(1, 0)	(1, 1)	(1, 2)	(2, 0)	(2, 1)	(2, 2)
Block	B1	B1	B2	B2	B3	B3	B4	B4	B5

Cache Hit / Miss and Locality

- Cache Miss vs Cache Hit
- Temporal Locality vs Spatial Locality
 - A program with good locality can significantly reduce miss rate and their penalty
- Block is the basic unit when transfer between memory
- Average Memory Access Time = Hit Time + Miss Rate * Miss Penalty

Cache Hit / Miss and Locality

- Cache Miss vs Cache Hit
- Temporal Locality vs Spatial Locality
 - A program with good locality can significantly reduce miss rate and their penalty
- Block is the basic unit when transfer between memory
- Average Memory Access Time = Hit Time + Miss Rate * Miss Penalty
 - Hit Time: Time to load data from cache if it is a cache hit

Cache Hit / Miss and Locality

- Cache Miss vs Cache Hit
- Temporal Locality vs Spatial Locality
 - A program with good locality can significantly reduce miss rate and their penalty
- Block is the basic unit when transfer between memory
- Average Memory Access Time = Hit Time + Miss Rate * Miss Penalty
 - Hit Time: Time to load data from cache if it is a cache hit
 - Miss Rate: Number of Miss / (Number of Hit + Number of Miss)

Cache Hit / Miss and Locality

- Cache Miss vs Cache Hit
- Temporal Locality vs Spatial Locality
 - A program with good locality can significantly reduce miss rate and their penalty
- Block is the basic unit when transfer between memory
- Average Memory Access Time = Hit Time + Miss Rate * Miss Penalty
 - Hit Time: Time to load data from cache if it is a cache hit
 - Miss Rate: Number of Miss / (Number of Hit + Number of Miss)
 - Miss Penalty: Extra Time to load data to the cache if it is a cache miss

Cache Hit / Miss and Locality

- Cache Miss vs Cache Hit
 - Temporal Locality vs Spatial Locality
 - A program with good locality can significantly reduce miss rate and their penalty
 - Block is the basic unit when transfer between memory
 - Average Memory Access Time = Hit Time + Miss Rate * Miss Penalty
 - Hit Time: Time to load data from cache if it is a cache hit
 - Miss Rate: Number of Miss / (Number of Hit + Number of Miss)
 - Miss Penalty: Extra Time to load data to the cache if it is a cache miss
- > **Hit Rate = 1 - Miss Rate**

Cache Hit / Miss and Locality

- Cache Miss vs Cache Hit
 - Temporal Locality vs Spatial Locality
 - A program with good locality can significantly reduce miss rate and their penalty
 - Block is the basic unit when transfer between memory
 - Average Memory Access Time = Hit Time + Miss Rate * Miss Penalty
 - Hit Time: Time to load data from cache if it is a cache hit
 - Miss Rate: Number of Miss / (Number of Hit + Number of Miss)
 - Miss Penalty: Extra Time to load data to the cache if it is a cache miss
- > Miss Time = Hit Time + Miss Penalty

Cache Hit / Miss and Locality

- Cache Miss vs Cache Hit
 - Temporal Locality vs Spatial Locality
 - A program with good locality can significantly reduce miss rate and their penalty
 - Block is the basic unit when transfer between memory
 - Average Memory Access Time = Hit Time + Miss Rate * Miss Penalty
 - Hit Time: Time to load data from cache if it is a cache hit
 - Miss Rate: Number of Miss / (Number of Hit + Number of Miss)
 - Miss Penalty: Extra Time to load data to the cache if it is a cache miss
- > **Memory Access Time = Average Memory Access Time * (Number of Hit + Number of Miss)**

Cache Hit / Miss and Locality

- Cache Miss vs Cache Hit
 - Temporal Locality vs Spatial Locality
 - A program with good locality can significantly reduce miss rate and their penalty
 - Block is the basic unit when transfer between memory
 - Average Memory Access Time = Hit Time + Miss Rate * Miss Penalty
 - Hit Time: Time to load data from cache if it is a cache hit
 - Miss Rate: Number of Miss / (Number of Hit + Number of Miss)
 - Miss Penalty: Extra Time to load data to the cache if it is a cache miss
- > $\text{Memory Access Time} = \text{Average Memory Access Time} * (\text{Number of Hit} + \text{Number of Miss})$
 $= \text{Hit Time} * \text{Number of Hit} + \text{Miss Time} * \text{Number of Miss}$

Cache Hit / Miss and Locality

- Cache Miss vs Cache Hit
 - Temporal Locality vs Spatial Locality
 - A program with good locality can significantly reduce miss rate and their penalty
 - Block is the basic unit when transfer between memory
 - Average Memory Access Time = Hit Time + Miss Rate * Miss Penalty
 - Hit Time: Time to load data from cache if it is a cache hit
 - Miss Rate: Number of Miss / (Number of Hit + Number of Miss)
 - Miss Penalty: Extra Time to load data to the cache if it is a cache miss
- > $\text{Memory Access Time} = \text{Average Memory Access Time} * (\text{Number of Hit} + \text{Number of Miss})$
 $= \text{Hit Time} * \text{Number of Hit} + \text{Miss Time} * \text{Number of Miss}$
 $= \text{Hit Time} * \text{Number of Hit} + (\text{Hit Time} + \text{Miss Penalty}) * \text{Number of Miss}$

Tag, set, and offset bits in memory address

In a 8-bit memory address system. Let $C[s][l][b]$ denote the value in cache in set s , line l , and byte b line l , all of which are zero-indexed, for instance $C[0][1][2]$ refers to the third byte in the second line of the first set. Let $S=2$ be the number of sets, $E=8$ be the number of lines per set, and $B=16$ be the number of bytes per block (the block being the data portion of a given line). Assume the tag field is in the most-significant bits [leftmost] and the block field is the least-significant bit [rightmost], What is the memory address (in hexadecimal) of that data in $C[0][2][7]$ with the tag bit 2?

Tag, set, and offset bits in memory address

In a 8-bit memory address system. Let $C[s][l][b]$ denote the value in cache in set s , line l , and byte b line l , all of which are zero-indexed, for instance $C[0][1][2]$ refers to the third byte in the second line of the first set. Let $S=2$ be the number of sets, $E=8$ be the number of lines per set, and $B=16$ be the number of bytes per block (the block being the data portion of a given line). Assume the tag field is in the most-significant bits [leftmost] and the block field is the least-significant bit [rightmost], What is the memory address (in hexadecimal) of that data in $C[0][2][7]$ with the tag bit 2?

Cache

S=0	16 bytes	16 bytes	16 bytes	16 bytes	16 bytes	16 bytes	16 bytes	16 bytes
S=1	16 bytes	16 bytes	16 bytes	16 bytes	16 bytes	16 bytes	16 bytes	16 bytes

Tag, set, and offset bits in memory address

In a 8-bit memory address system. Let $C[s][l][b]$ denote the value in cache in set s , line l , and byte b line l , all of which are zero-indexed, for instance $C[0][1][2]$ refers to the third byte in the second line of the first set. Let $S=2$ be the number of sets, $E=8$ be the number of lines per set, and $B=16$ be the number of bytes per block (the block being the data portion of a given line). Assume the tag field is in the most-significant bits [leftmost] and the block field is the least-significant bit [rightmost], What is the memory address (in hexadecimal) of that data in $C[0][2][7]$ with the tag bit 2?

8 bits		
	Set=1	

Tag, set, and offset bits in memory address

In a 8-bit memory address system. Let $C[s][l][b]$ denote the value in cache in set s , line l , and byte b line l , all of which are zero-indexed, for instance $C[0][1][2]$ refers to the third byte in the second line of the first set. Let $S=2$ be the number of sets, $E=8$ be the number of lines per set, and $B=16$ be the number of bytes per block (the block being the data portion of a given line). Assume the tag field is in the most-significant bits [leftmost] and the block field is the least-significant bit [rightmost], What is the memory address (in hexadecimal) of that data in $C[0][2][7]$ with the tag bit 2?

8 bits		
	Set=1	Offset=4

Tag, set, and offset bits in memory address

In a 8-bit memory address system. Let $C[s][l][b]$ denote the value in cache in set s , line l , and byte b line l , all of which are zero-indexed, for instance $C[0][1][2]$ refers to the third byte in the second line of the first set. Let $S=2$ be the number of sets, $E=8$ be the number of lines per set, and $B=16$ be the number of bytes per block (the block being the data portion of a given line). Assume the tag field is in the most-significant bits [leftmost] and the block field is the least-significant bit [rightmost], What is the memory address (in hexadecimal) of that data in $C[0][2][7]$ with the tag bit 2?

8 bits		
Tag=8-1-4=3	Set=1	Offset=4

Tag, set, and offset bits in memory address

In a 8-bit memory address system. Let $C[s][l][b]$ denote the value in cache in set s , line l , and byte b line l , all of which are zero-indexed, for instance $C[0][1][2]$ refers to the third byte in the second line of the first set. Let $S=2$ be the number of sets, $E=8$ be the number of lines per set, and $B=16$ be the number of bytes per block (the block being the data portion of a given line). Assume the tag field is in the most-significant bits [leftmost] and the block field is the least-significant bit [rightmost], What is the memory address (in hexadecimal) of that data in $C[0][2][7]$ with the tag bit 2?

8 bits		
Tag=3	Set=1	Offset=4
010	0	0111

Answer is 0100 0111 in binary, which is 47 in hexadecimal

Cache Placement Policy

Consider a 2-way set associative cache running LRU replacement algorithm that services a memory space of 32 byte locations. 5-bit memory addresses are divided into the format of [ttssb]. In each set, the top row is most recently used and bottom row is least recently used. Indicate whether each of the following memory accesses was a hit or a miss. Noticed that all given addresses are independent, so the state of the cache is not changed after accessing each of the given addresses.

11010

Set	Valid	Tag	Block
0	1	10	
	1	11	
1	1	11	
	0	00	
2	1	00	
	0	11	
3	0	01	
	0	00	

Cache Placement Policy

Consider a 2-way set associative cache running LRU replacement algorithm that services a memory space of 32 byte locations. 5-bit memory addresses are divided into the format of [ttssb]. In each set, the top row is most recently used and bottom row is least recently used. Indicate whether each of the following memory accesses was a hit or a miss. Noticed that all given addresses are independent, so the state of the cache is not changed after accessing each of the given addresses.

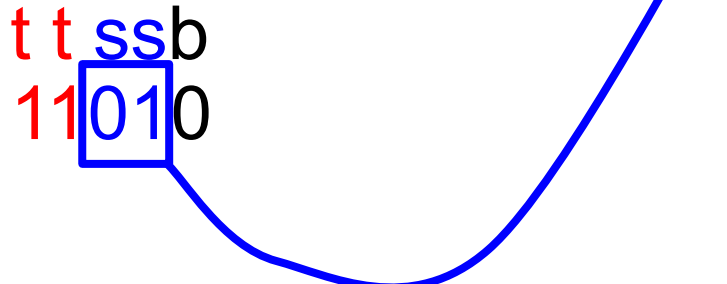
ttssb
11010

Set	Valid	Tag	Block
0	1	10	
	1	11	
1	1	11	
	0	00	
2	1	00	
	0	11	
3	0	01	
	0	00	

Cache Placement Policy

Consider a 2-way set associative cache running LRU replacement algorithm that services a memory space of 32 byte locations. 5-bit memory addresses are divided into the format of [ttssb]. In each set, the top row is most recently used and bottom row is least recently used. Indicate whether each of the following memory accesses was a hit or a miss. Noticed that all given addresses are independent, so the state of the cache is not changed after accessing each of the given addresses.

ttssb
11010

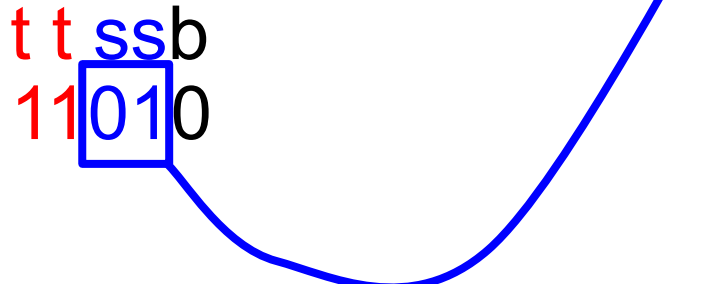


Set	Valid	Tag	Block
0	1	10	
	1	11	
1	1	11	
	0	00	
2	1	00	
	0	11	
3	0	01	
	0	00	

Cache Placement Policy

Consider a 2-way set associative cache running LRU replacement algorithm that services a memory space of 32 byte locations. 5-bit memory addresses are divided into the format of [ttssb]. In each set, the top row is most recently used and bottom row is least recently used. Indicate whether each of the following memory accesses was a hit or a miss. Noticed that all given addresses are independent, so the state of the cache is not changed after accessing each of the given addresses.

ttssb
11010



Set	Valid	Tag	Block
0	1	10	
	1	11	
1	1	11	
	0	00	
2	1	00	
	0	11	
3	0	01	
	0	00	

Cache Placement Policy


Consider a 2-way set associative cache running LRU replacement algorithm that services a memory space of 32 byte locations. 5-bit memory addresses are divided into the format of [ttssb]. In each set, the top row is most recently used and bottom row is least recently used. Indicate whether each of the following memory accesses was a hit or a miss. Noticed that all given addresses are independent, so the state of the cache is not changed after accessing each of the given addresses.

ttssb
11010

Set	Valid	Tag	Block
0	1	10	
	1	11	
1	1	11	
	0	00	
2	1	00	
	0	11	
3	0	01	
	0	00	

Cache Placement Policy

Consider a 2-way set associative cache running LRU replacement algorithm that services a memory space of 32 byte locations. 5-bit memory addresses are divided into the format of [ttssb]. In each set, the top row is most recently used and bottom row is least recently used. Indicate whether each of the following memory accesses was a hit or a miss. Noticed that all given addresses are independent, so the state of the cache is not changed after accessing each of the given addresses.



HIT, load the first byte.

Set	Valid	Tag	Block
0	1	10	
	1	11	
1	1	11	
	0	00	
2	1	00	
	0	11	
3	0	01	
	0	00	

Cache Placement Policy

Consider a 2-way set associative cache running LRU replacement algorithm that services a memory space of 32 byte locations. 5-bit memory addresses are divided into the format of [ttssb]. In each set, the top row is most recently used and bottom row is least recently used. Indicate whether each of the following memory accesses was a hit or a miss. Noticed that all given addresses are independent, so the state of the cache is not changed after accessing each of the given addresses.

01110

Set	Valid	Tag	Block
0	1	10	
	1	11	
1	1	11	
	0	00	
2	1	00	
	0	11	
3	0	01	
	0	00	

Cache Placement Policy

Consider a 2-way set associative cache running LRU replacement algorithm that services a memory space of 32 byte locations. 5-bit memory addresses are divided into the format of [ttssb]. In each set, the top row is most recently used and bottom row is least recently used. Indicate whether each of the following memory accesses was a hit or a miss. Noticed that all given addresses are independent, so the state of the cache is not changed after accessing each of the given addresses.

ttssb
01110

Set	Valid	Tag	Block
0	1	10	
	1	11	
1	1	11	
	0	00	
2	1	00	
	0	11	
3	0	01	
	0	00	

Cache Placement Policy

Consider a 2-way set associative cache running LRU replacement algorithm that services a memory space of 32 byte locations. 5-bit memory addresses are divided into the format of [ttssb]. In each set, the top row is most recently used and bottom row is least recently used. Indicate whether each of the following memory accesses was a hit or a miss. Noticed that all given addresses are independent, so the state of the cache is not changed after accessing each of the given addresses.

ttssb
01110

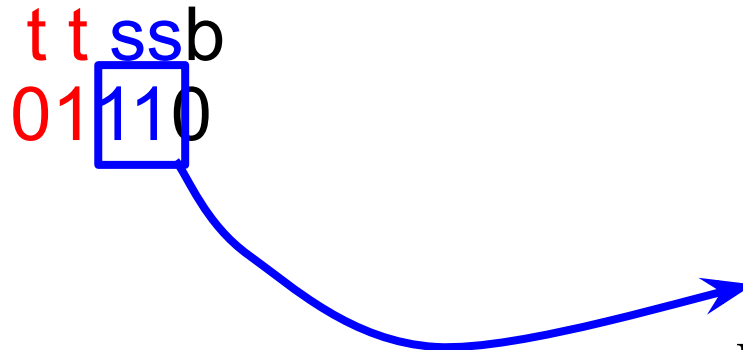


Set	Valid	Tag	Block
0	1	10	
	1	11	
1	1	11	
	0	00	
2	1	00	
	0	11	
3	0	01	
	0	00	

Cache Placement Policy

Consider a 2-way set associative cache running LRU replacement algorithm that services a memory space of 32 byte locations. 5-bit memory addresses are divided into the format of [ttssb]. In each set, the top row is most recently used and bottom row is least recently used. Indicate whether each of the following memory accesses was a hit or a miss. Noticed that all given addresses are independent, so the state of the cache is not changed after accessing each of the given addresses.

ttssb
01110



Set	Valid	Tag	Block
0	1	10	
	1	11	
1	1	11	
	0	00	
2	1	00	
	0	11	
3	0	01	
	0	00	

Cache Placement Policy

Consider a 2-way set associative cache running LRU replacement algorithm that services a memory space of 32 byte locations. 5-bit memory addresses are divided into the format of [ttssb]. In each set, the top row is most recently used and bottom row is least recently used. Indicate whether each of the following memory accesses was a hit or a miss. Noticed that all given addresses are independent, so the state of the cache is not changed after accessing each of the given addresses.

ttssb
01110



Set	Valid	Tag	Block
0	1	10	
	1	11	
1	1	11	
	0	00	
2	1	00	
	0	11	
3	0	01	
	0	00	

Cache Placement Policy

Consider a 2-way set associative cache running LRU replacement algorithm that services a memory space of 32 byte locations. 5-bit memory addresses are divided into the format of [ttssb]. In each set, the top row is most recently used and bottom row is least recently used. Indicate whether each of the following memory accesses was a hit or a miss. Noticed that all given addresses are independent, so the state of the cache is not changed after accessing each of the given addresses.

01110

MISS, check the next level.

Set	Valid	Tag	Block
0	1	10	
	1	11	
1	1	11	
	0	00	
2	1	00	
	0	11	
3	0	01	
	0	00	

Cache Replacement Policy

$A = [4, 1, 4, 3, 4, 1, 2, 3, 3, 4, 3, 1, 3, 2, 1, 3]$

Suppose A, a list of address requests. We execute those requests starting with the request at index 0 in the list. We are using the **Least Recently Used** replacement policy to cache those request. The capacity of the cache is 3 addresses.

Suppose A, a list of address requests. We execute those requests starting with the request at index 0 in the list. We are using the **Least Frequently Used** replacement policy to cache those request. The capacity of the cache is 3 addresses. If there is a tie, the smallest number is removed.

Cache Replacement Policy

$A = [4, 1, 4, 3, 4, 1, 2, 3, 3, 4, 3, 1, 3, 2, 1, 3]$

Suppose A, a list of address requests. We execute those requests starting with the request at index 0 in the list. We are using the **Least Recently Used** replacement policy to cache those request. The capacity of the cache is 3 addresses.

A: 1, 2, 3 are the final data in the cache after visiting all the data since they are the most recent 3 elements

Suppose A, a list of address requests. We execute those requests starting with the request at index 0 in the list. We are using the **Least Frequently Used** replacement policy to cache those request. The capacity of the cache is 3 addresses. If there is a tie, the smallest number is removed.

Cache Replacement Policy

$A = [4, 1, 4, 3, 4, 1, 2, 3, 3, 4, 3, 1, 3, 2, 1, 3]$

Suppose A, a list of address requests. We execute those requests starting with the request at index 0 in the list. We are using the **Least Recently Used** replacement policy to cache those request. The capacity of the cache is 3 addresses.

A: 1, 2, 3 are the final data in the cache after visiting all the data since they are the most recent 3 elements

Suppose A, a list of address requests. We execute those requests starting with the request at index 0 in the list. We are using the **Least Frequently Used** replacement policy to cache those request. The capacity of the cache is 3 addresses. If there is a tie, the smallest number is removed.

A: 1 3 4 are the final data in the cache after visiting all the data since the frequency of them are 1 5 4. Notice that the frequency is reset to 0 when the data is evicted from the cache during the replacement process.


Systems Fundamentals

Things you should know from this chapter:

- What are system calls and library calls?
- How to do error handling for system calls? [errno]
- What are 3 OS abstraction?
- What are the 3 table used to manipulate files?
- System calls used for manipulating files:
 - open
 - read
 - write
 - close
 - lseek

System Calls vs Library Calls

Figure taken from
lecture slides:



Are system calls always
faster than their library
counterparts?

False

ERRNO

Taken from lecture slides:

- When a system call fails, it sets the global integer variable `errno` to a positive value that identifies the specific error
- Including the `<errno.h>` header file provides a declaration of `errno`, as well as a set of constants for the various error numbers
- The section headed `ERRORS` in each manual page provides a list of possible `errno` values that can be returned by each system call

ERRNO

Taken from lecture slides:

- When a system call fails, it sets the global integer variable `errno` to a positive value that identifies the specific error
- Including the `<errno.h>` header file provides a declaration of `errno`, as well as a set of constants for the various error numbers
- The section headed `ERRORS` in each manual page provides a list of possible `errno` values that can be returned by each system call

`errno` gets reset to 0, everytime a system call is successful?

ERRNO

Taken from lecture slides:

- When a system call fails, it sets the global integer variable `errno` to a positive value that identifies the specific error
- Including the `<errno.h>` header file provides a declaration of `errno`, as well as a set of constants for the various error numbers
- The section headed `ERRORS` in each manual page provides a list of possible `errno` values that can be returned by each system call

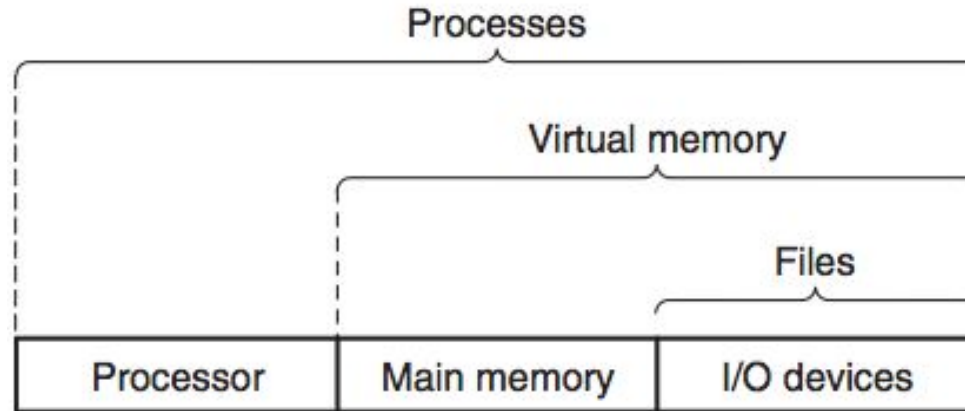
`errno` gets reset to 0, everytime a system call is successful?

False

What are the 3 OS Abstractions?


What are the 3 OS Abstractions?

Figure taken from lecture slides:



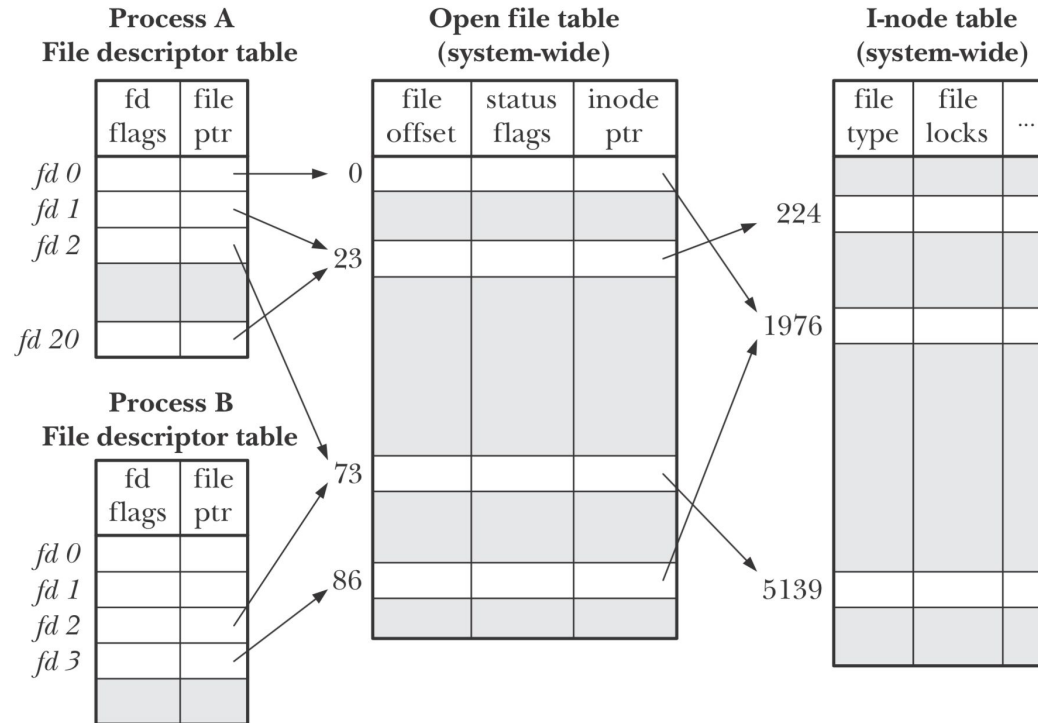
What are the 3 OS Abstractions?

Figure taken from lecture slides:



Now recap from the worksheet 6!

Three tables behind the files system



What happens when I run this code?

Process A
File Descriptor Table

fd flags	file ptr
....	...
...	...
...	...

Open File Table
(System Wide)

file offset	status flag	inode pointer
.	.	.
.	.	.
.	.	.

```
int fd1 = open("file1.txt", O_RDWR);
int fd2 = open("file2.txt", O_RDWR);
int fd3 = dup(fd1)
int fd4 = open("file1.txt", O_RDWR);
```

```
write(fd1, "CS230 ", 6);
write(fd3, "Nikko", 5);
write(fd4, "Joe C", 5);
```

```
lseek(fd3, -3, SEEK_CUR);
lseek(fd4, 0, SEEK_END);
```

file1.txt:

ABCD

What are the first 3 entry in file descriptor table?

Figure taken from lecture slides:

What will be the final output of this code?

File descriptor	Purpose	POSIX name	stdio stream
0	standard input	STDIN_FILENO	<i>stdin</i>
1	standard output	STDOUT_FILENO	<i>stdout</i>
2	standard error	STDERR_FILENO	<i>stderr</i>

What happens when I run this code?

Process A
File Descriptor Table

fd flags	file ptr
....	...
...	...
...	...
RD,WR	20

Open File Table
(System Wide)

file offset	status flag	inode pointer
.	.	.
.	.	.
.	.	.
0	...	100

```
int fd1 = open("file1.txt", O_RDWR);  
int fd2 = open("file2.txt", O_RDWR);  
int fd3 = dup(fd1)  
int fd4 = open("file1.txt", O_RDWR);
```

```
write(fd1, "CS230 ", 6);  
write(fd3, "Nikko", 5);  
write(fd4, "Joe C", 5);
```

```
lseek(fd3, -3, SEEK_CUR);  
lseek(fd4, 0, SEEK_END);
```

file1.txt:

ABCD

What happens when I run this code?

```
int fd1 = open("file1.txt", O_RDWR);
int fd2 = open("file2.txt", O_RDWR);
int fd3 = dup(fd1)
int fd4 = open("file1.txt", O_RDWR);
```

```
write(fd1, "CS230 ", 6);
write(fd3, "Nikko", 5);
write(fd4, "Joe C", 5);
```

```
lseek(fd3, -3, SEEK_CUR);
lseek(fd4, 0, SEEK_END);
```

Process A
File Descriptor Table

fd flags	file ptr
....	...
...	...
...	...
RD,WR	20
RD,WR	21

Open File Table
(System Wide)

file offset	status flag	inode pointer
.	.	.
.	.	.
.	.	.
0	...	100
0	...	250

file1.txt:

What happens when I run this code?

Process A
File Descriptor Table

fd flags	file ptr
....	...
...	...
...	...
RD,WR	20
RD,WR	21
RD,WR	20

Open File Table
(System Wide)

file offset	status flag	inode pointer
.	.	.
.	.	.
.	.	.
0	...	100
0	...	250

```
int fd1 = open("file1.txt", O_RDWR);  
int fd2 = open("file2.txt", O_RDWR);  
int fd3 = dup(fd1);  
int fd4 = open("file1.txt", O_RDWR);
```

```
write(fd1, "CS230 ", 6);  
write(fd3, "Nikko", 5);  
write(fd4, "Joe C", 5);
```

```
lseek(fd3, -3, SEEK_CUR);  
lseek(fd4, 0, SEEK_END);
```

file1.txt:

ABCD

What happens when I run this code?

Process A
File Descriptor Table

fd flags	file ptr
....	...
...	...
...	...
RD,WR	20
RD,WR	21
RD,WR	20
RD,WR	22

Open File Table
(System Wide)

file offset	status flag	inode pointer
.	.	.
.	.	.
.	.	.
0	...	100
0	...	250
0	...	100

```
int fd1 = open("file1.txt", O_RDWR);
int fd2 = open("file2.txt", O_RDWR);
int fd3 = dup(fd1);
int fd4 = open("file1.txt", O_RDWR);
```

```
write(fd1, "CS230 ", 6);
write(fd3, "Nikko", 5);
write(fd4, "Joe C", 5);
```

```
lseek(fd3, -3, SEEK_CUR);
lseek(fd4, 0, SEEK_END);
```

file1.txt:



What happens when I run this code?

```
int fd1 = open("file1.txt", O_RDWR);
int fd2 = open("file2.txt", O_RDWR);
int fd3 = dup(fd1);
int fd4 = open("file1.txt", O_RDWR);
```

```
write(fd1, "CS230 ", 6);
write(fd3, "Nikko", 5);
write(fd4, "Joe C", 5);
```

```
lseek(fd3, -3, SEEK_CUR);
lseek(fd4, 0, SEEK_END);
```

file1.txt:

|||ABCD => |CS230 |||


Aadam Lokhandwala, Joe Chiu

Process A
File Descriptor Table

fd flags	file ptr
....	...
...	...
...	...
RD,WR	20
RD,WR	21
RD,WR	20
RD,WR	22

Open File Table
(System Wide)

file offset	status flag	inode pointer
.	.	.
.	.	.
.	.	.
0 => 6	...	100
0	...	250
0	...	100



What happens when I run this code?

Process A
File Descriptor Table

fd flags	file ptr
....	...
...	...
...	...
RD,WR	20
RD,WR	21
RD,WR	20
RD,WR	22

Open File Table
(System Wide)

file offset	status flag	inode pointer
.	.	.
.	.	.
.	.	.
6 => 11	...	100
0	...	250
0	...	100

```
int fd1 = open("file1.txt", O_RDWR);
int fd2 = open("file2.txt", O_RDWR);
int fd3 = dup(fd1);
int fd4 = open("file1.txt", O_RDWR);
```

```
write(fd1, "CS230 ", 6);
write(fd3, "Nikko", 5);
write(fd4, "Joe C", 5);
```

```
lseek(fd3, -3, SEEK_CUR);
lseek(fd4, 0, SEEK_END);
```

file1.txt:

|CS230 || => |CS230 Nikko||

Aadam Lokhandwala, Joe Chiu

UMassAmherst

College of Information
& Computer Sciences

What happens when I run this code?

Process A
File Descriptor Table

fd flags	file ptr
....	...
...	...
...	...
RD,WR	20
RD,WR	21
RD,WR	20
RD,WR	22

Open File Table
(System Wide)

file offset	status flag	inode pointer
.	.	.
.	.	.
.	.	.
11	...	100
0	...	250
0 => 5	...	100

```
int fd1 = open("file1.txt", O_RDWR);
int fd2 = open("file2.txt", O_RDWR);
int fd3 = dup(fd1);
int fd4 = open("file1.txt", O_RDWR);
```

```
write(fd1, "CS230 ", 6);
write(fd3, "Nikko", 5);
write(fd4, "Joe C", 5);
```

```
lseek(fd3, -3, SEEK_CUR);
lseek(fd4, 0, SEEK_END);
```

file1.txt:

|CS230 Nikko|| => Joe C| Nikko||

Aadam Lokhandwala, Joe Chiu

What happens when I run this code?

Process A
File Descriptor Table

fd flags	file ptr
....	...
...	...
...	...
RD,WR	20
RD,WR	21
RD,WR	20
RD,WR	22

Open File Table
(System Wide)

file offset	status flag	inode
.	.	.
.	.	.
.	.	.
11 - 3 => 8	...	100
0	...	250
5	...	100

```
int fd1 = open("file1.txt", O_RDWR);
int fd2 = open("file2.txt", O_RDWR);
int fd3 = dup(fd1);
int fd4 = open("file1.txt", O_RDWR);
```

```
write(fd1, "CS230 ", 6);
write(fd3, "Nikko", 5);
write(fd4, "Joe C", 5);
```

```
lseek(fd3, -3, SEEK_CUR);
lseek(fd4, 0, SEEK_END);
```

file1.txt:

Joe C | Nikko | => Joe C | Ni | kko

Aadam Lokhandwala, Joe Chiu

What happens when I run this code?

Process A
File Descriptor Table

fd flags	file ptr
....	...
...	...
...	...
RD,WR	20
RD,WR	21
RD,WR	20
RD,WR	22

Open File Table
(System Wide)

file offset	status flag	inode
.	.	.
.	.	.
.	.	.
8	...	100
0	...	250
5 => 11	...	100

```
int fd1 = open("file1.txt", O_RDWR);
int fd2 = open("file2.txt", O_RDWR);
int fd3 = dup(fd1);
int fd4 = open("file1.txt", O_RDWR);
```

```
write(fd1, "CS230 ", 6);
write(fd3, "Nikko", 5);
write(fd4, "Joe C", 5);
```

```
lseek(fd3, -3, SEEK_CUR);
lseek(fd4, 0, SEEK_END);
```

file1.txt:

Joe C | Ni | kko => Joe C Ni | kko |

Aadam Lokhandwala, Joe Chiu

What happens when I run this code?

Process A
File Descriptor Table

fd flags	file ptr
....	...
...	...
...	...
RD,WR	20
RD,WR	21
RD,WR	20
RD,WR	22

Open File Table
(System Wide)

file offset	status flag	inode
.	.	.
.	.	.
.	.	.
8	...	100
0	...	250
11	...	100

```
int fd1 = open("file1.txt", O_RDWR);
int fd2 = open("file2.txt", O_RDWR);
int fd3 = dup(fd1);
int fd4 = open("file1.txt", O_RDWR);
```

```
write(fd1, "CS230 ", 6);
write(fd3, "Nikko", 5);
write(fd4, "Joe C", 5);
```

```
lseek(fd3, -3, SEEK_CUR);
lseek(fd4, 0, SEEK_END);
```

file1.txt:

Joe C Ni||kko|

Aadam Lokhandwala, Joe Chiu

UMassAmherst

College of Information
& Computer Sciences

Virtual Memory

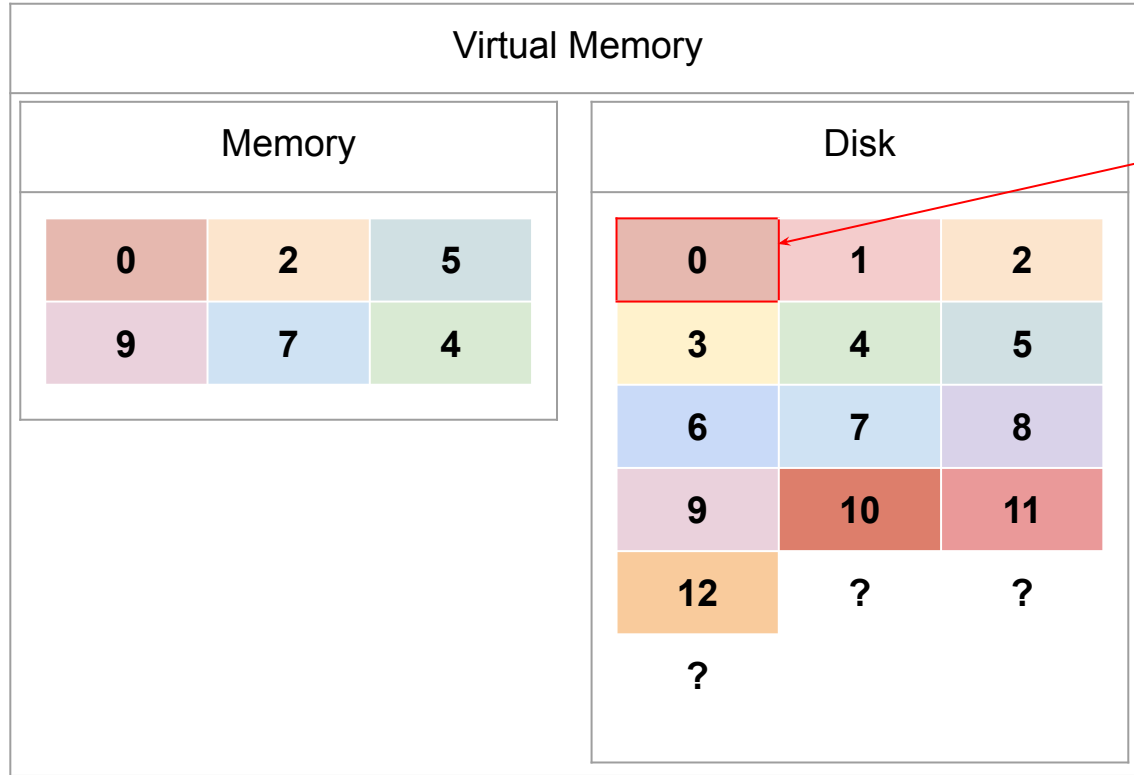
Things you should know from this chapter:

- How does Page Table works
- How to identify a page is Allocated/ Unallocated, Cached/Uncached
- What happens when page is not cached? (Page Fault)
- Address Translation (How to convert VA to PA)

Virtual Memory

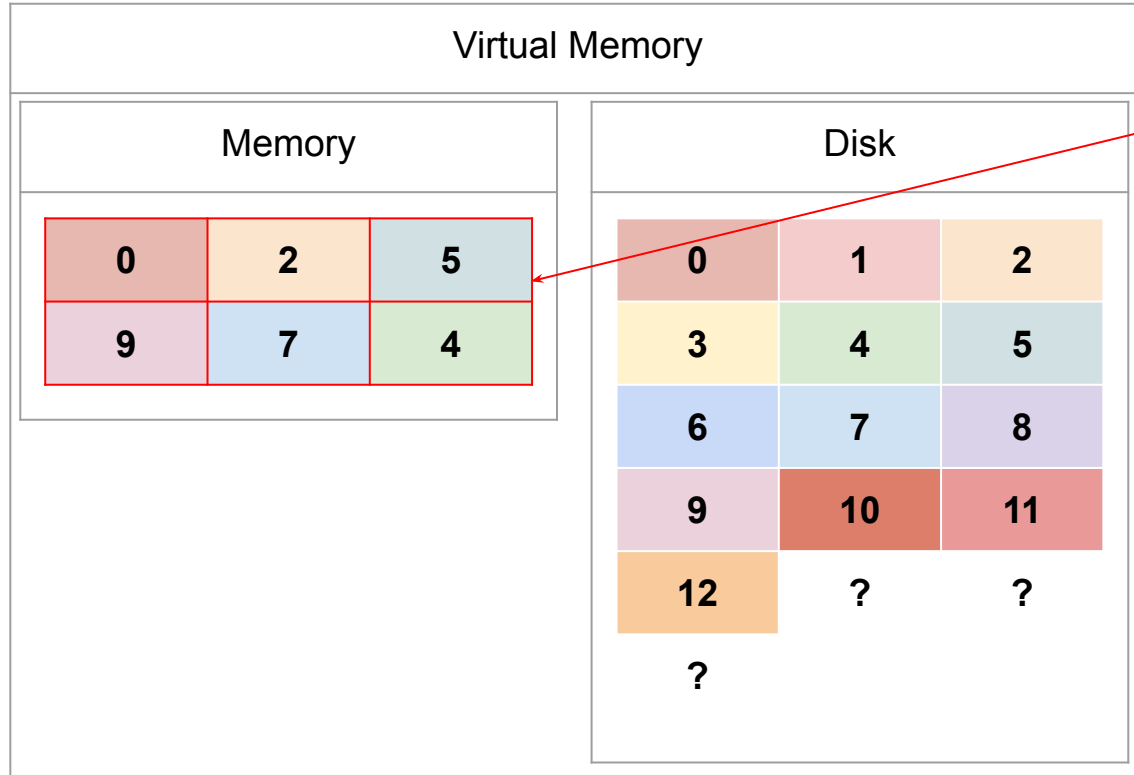
Virtual Memory					
Memory			Disk		
0	2	5	0	1	2
9	7	4	3	4	5
			6	7	8
			9	10	11
			12	?	?
			?		

Virtual Memory



Page

Virtual Memory

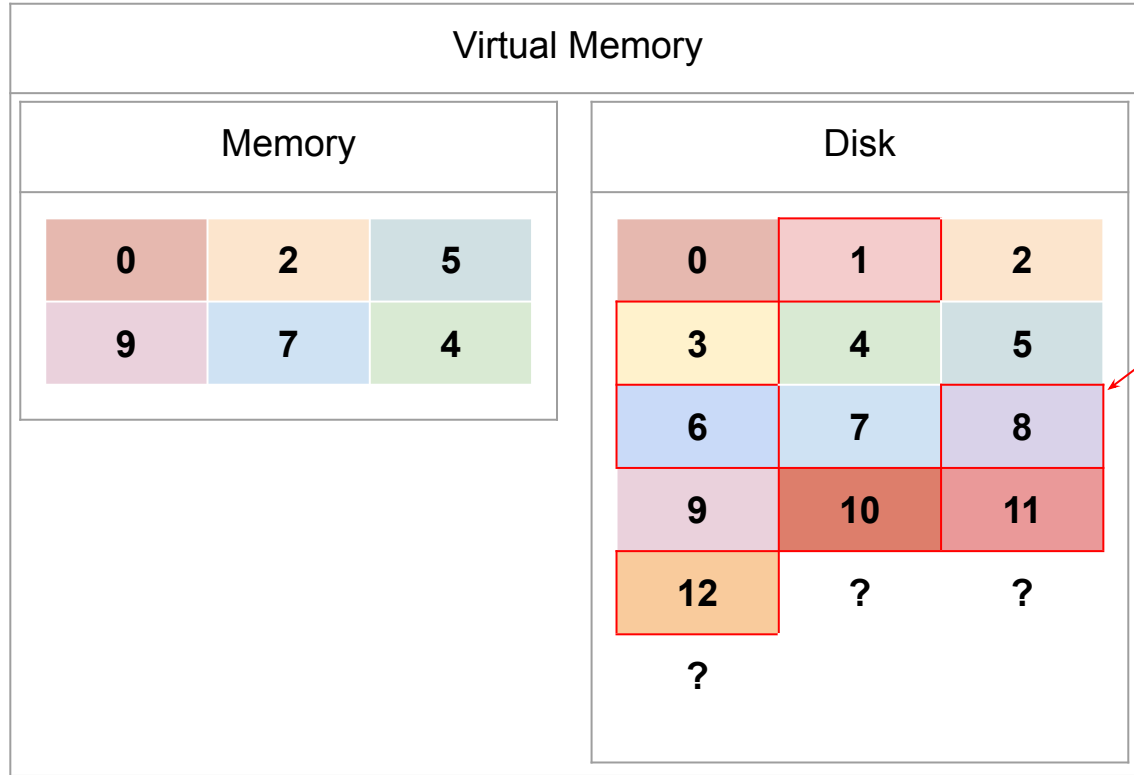


Allocated & Cached

Allocated & Un-Cached

Unallocated

Virtual Memory




Allocated & Cached

Allocated & Un-Cached

Unallocated

Virtual Memory



Allocated & Cached

Allocated & Un-Cached


Unallocated

Virtual Memory

Virtual Memory					
Memory			Disk		
0	2	5	0	1	2
9	7	4	3	4	5
			6	7	8
			9	10	11
			12	?	?
			?		

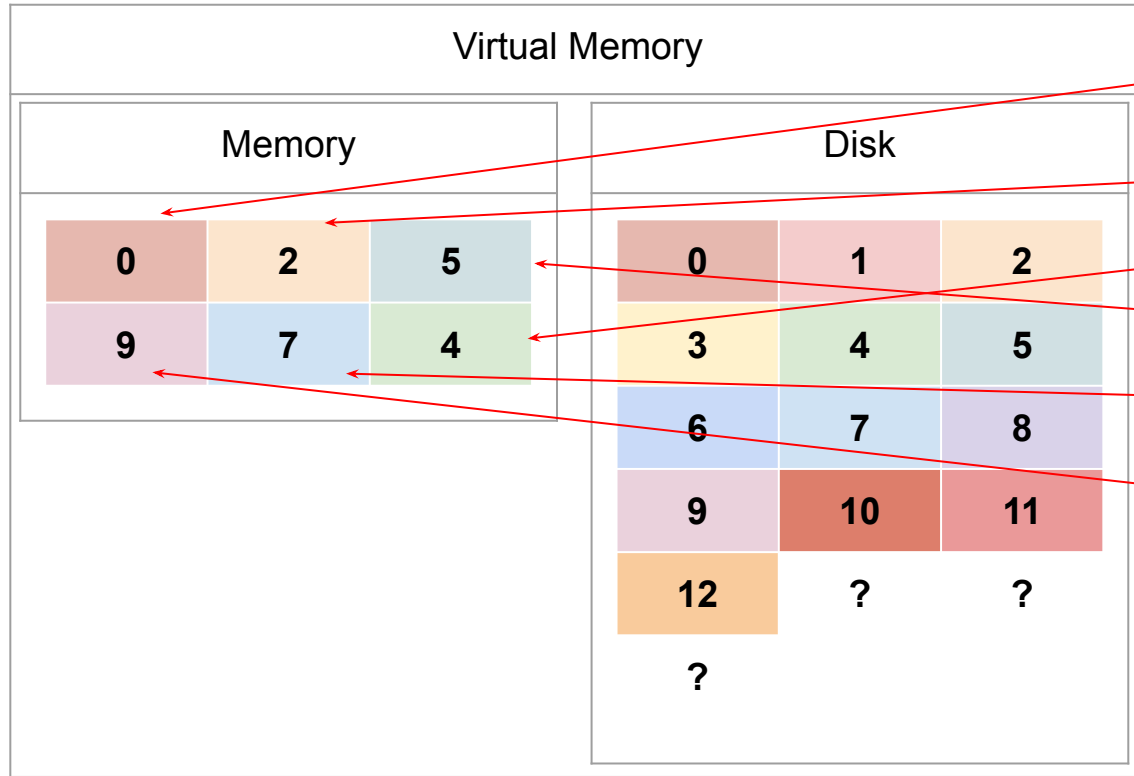
Page Table	
1	0000
0	0001
1	0001
0	0011
1	0010
1	0101
0	0110
1	0100
0	1000
1	0011
0	1010
0	1011
0	1100
0	Null
0	Null
0	Null

Virtual Memory



Page Table		
1	0000	Page Table Entry (PTE)
0	0001	
1	0001	
0	0011	
1	0010	
1	0101	Valid Bit
0	0110	
1	0100	
0	1000	
1	0011	Address
0	1010	
0	1011	
0	1100	
0	Null	
0	Null	
0	Null	

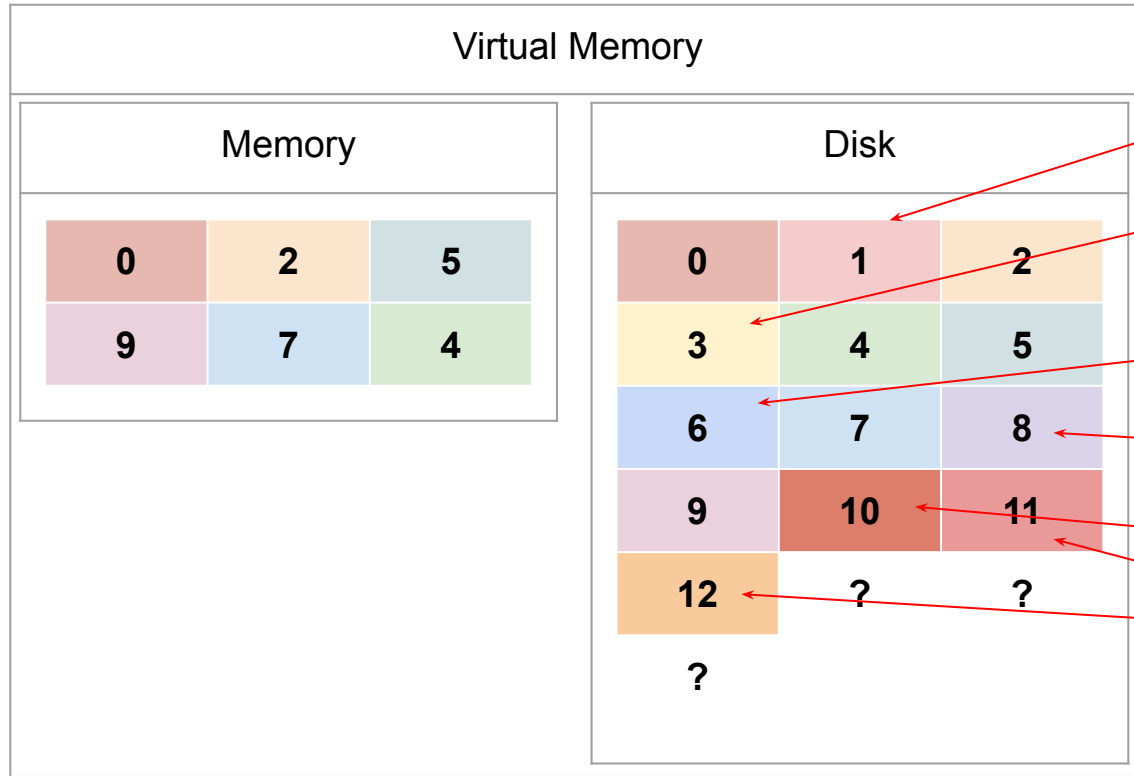
Virtual Memory



Page Table	
1	0000
0	0001
1	0001
0	0011
1	0010
1	0101
0	0110
1	0100
0	1000
1	0011
0	1010
0	1011
0	1100
0	Null
0	Null
0	Null

When **valid bit is 1**, it means the page is in memory. Hence, it's **Cached**

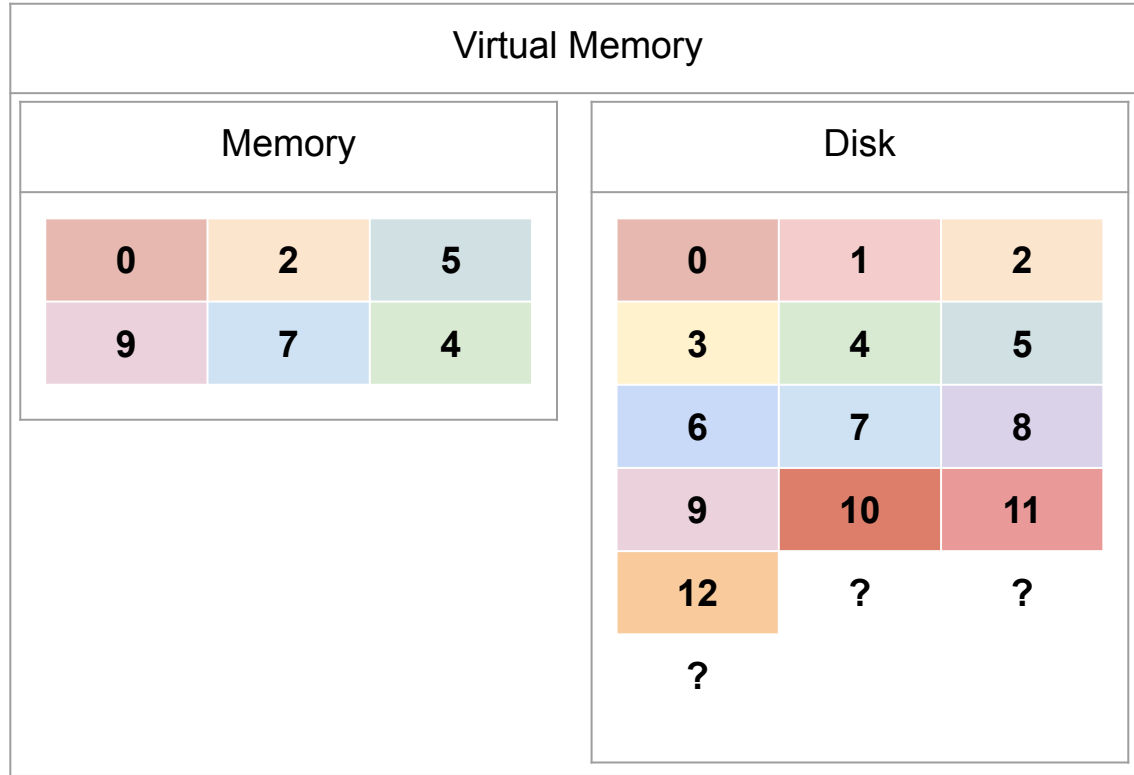
Virtual Memory



Page Table	
1	0000
0	0001
1	0001
0	0011
1	0010
1	0101
0	0110
1	0100
0	1000
1	0011
0	1010
0	1011
0	1100
0	Null
0	Null
0	Null

When **valid bit is 0** and **address is not null**, it means the page is not in memory. Hence, it's **Un-Cached**

Virtual Memory



Page Table	
1	0000
0	0001
1	0001
0	0011
1	0010
1	0101
0	0110
1	0100
0	1000
1	0011
0	1010
0	1011
0	1100
0	Null
0	Null
0	Null

When **valid bit is 0** and **address is null**, it means the page is not allocated yet. Hence, it's **Un-Allocated**

Address Translation

Page Table	
1	0000
0	0001
1	0001
0	0011
1	0010
1	0101
0	0110
1	0100
0	1000
1	0011
0	1010
0	1011
0	1100
0	Null
0	Null
0	Null

How can we convert Virtual Address to Physical Address?

First let's see how VA and PA looks like:

- Virtual Address:

Virtual Page Number(VPN) (n-p bits)	Page Offset(PO) (p bits)
--	-----------------------------

- Physical Address:

Physical Page Number(PPN) (m-p bits)	Page Offset(PO) (p bits)
---	-----------------------------

Address Translation

Page Table	
1	0000
0	0001
1	0001
0	0011
1	0010
1	0101
0	0110
1	0100
0	1000
1	0011
0	1010
0	1011
0	1100
0	Null
0	Null
0	Null

How can we convert Virtual Address to Physical Address?

First let's see how VA and PA looks like:

- Virtual Address:

Virtual Page Number(VPN) (n-p bits)	Page Offset(PO) (p bits)
--	-----------------------------

- Physical Address:

Physical Page Number(PPN) (m-p bits)	Page Offset(PO) (p bits)
---	-----------------------------

Which piece of hardware is responsible for converting VA to PA?

Address Translation

Page Table	
1	0000
0	0001
1	0001
0	0011
1	0010
1	0101
0	0110
1	0100
0	1000
1	0011
0	1010
0	1011
0	1100
0	Null
0	Null
0	Null

How can we convert Virtual Address to Physical Address?

First let's see how VA and PA looks like:

- Virtual Address:

Virtual Page Number(VPN) (n-p bits)	Page Offset(PO) (p bits)
--	-----------------------------

- Physical Address:

Physical Page Number(PPN) (m-p bits)	Page Offset(PO) (p bits)
---	-----------------------------

Which piece of hardware is responsible for converting VA to PA?
Memory Management Unit (MMU)

Address Translation

Page Table	
1	0000
0	0001
1	0001
0	0011
1	0010
1	0101
0	0110
1	0100
0	1000
1	0011
0	1010
0	1011
0	1100
0	Null
0	Null
0	Null

How can we convert Virtual Address to Physical Address?

First let's see how VA and PA looks like:

- Virtual Address:

Virtual Page Number(VPN) (n-p bits)	Page Offset(PO) (p bits)
--	-----------------------------

- Physical Address:

Physical Page Number(PPN) (m-p bits)	Page Offset(PO) (p bits)
---	-----------------------------

Which piece of hardware is responsible for converting VA to PA?
Memory Management Unit (MMU)

Can you convert from PA to VA?

Address Translation

Page Table	
1	0000
0	0001
1	0001
0	0011
1	0010
1	0101
0	0110
1	0100
0	1000
1	0011
0	1010
0	1011
0	1100
0	Null
0	Null
0	Null

How can we convert Virtual Address to Physical Address?

First let's see how VA and PA looks like:

- Virtual Address:

Virtual Page Number(VPN) (n-p bits)	Page Offset(PO) (p bits)
--	-----------------------------

- Physical Address:

Physical Page Number(PPN) (m-p bits)	Page Offset(PO) (p bits)
---	-----------------------------

Which piece of hardware is responsible for converting VA to PA?
Memory Management Unit (MMU)

Can you convert from PA to VA?

No, Address translation only works one way.

Address Translation

Page Table	
1	0
0	1
1	1
0	11
1	10
1	101
0	110
1	100
0	1000
1	11
0	1010
0	1011
0	1100
0	Null
0	Null
0	Null

How can we convert Virtual Address to Physical Address?

First let's see how VA and PA looks like:

- Virtual Address:

Virtual Page Number(VPN) (n-p bits)	Page Offset(PO) (p bits)
--	-----------------------------

- Physical Address:

Physical Page Number(PPN) (m-p bits)	Page Offset(PO) (p bits)
---	-----------------------------

How many bits will VA and PA require? Let's say we are given $M = 128$, and virtual memory uses 6 bits.

Address Translation

Page Table	
1	0
0	1
1	1
0	11
1	10
1	101
0	110
1	100
0	1000
1	11
0	1010
0	1011
0	1100
0	Null
0	Null
0	Null

How can we convert Virtual Address to Physical Address?

First let's see how VA and PA looks like:

- Virtual Address:

Virtual Page Number(VPN) (n-p bits)	Page Offset(PO) (p bits)
--	-----------------------------

- Physical Address:

Physical Page Number(PPN) (m-p bits)	Page Offset(PO) (p bits)
---	-----------------------------

How many bits will VA and PA require? Let's say we are given $M = 128$, and virtual memory uses 6 bits.

$$\text{As } M = 2^m \Rightarrow 128 = 2^7 \Rightarrow m = 7$$

$$\text{And as virtual memory uses 6 bits} \Rightarrow n = 6$$

Address Translation

Page Table	
1	0
0	1
1	1
0	11
1	10
1	101
0	110
1	100
0	1000
1	11
0	1010
0	1011
0	1100
0	Null
0	Null
0	Null

How can we convert Virtual Address to Physical Address?

First let's see how VA and PA looks like:

- Virtual Address:

Virtual Page Number(VPN) (6-p bits)	Page Offset(PO) (p bits)
--	-----------------------------

- Physical Address:

Physical Page Number(PPN) (7-p bits)	Page Offset(PO) (p bits)
---	-----------------------------

How many bits will VA and PA require? Let's say we are given $M = 128$, and virtual memory uses 6 bits.

$$\text{As } M = 2^m \Rightarrow 128 = 2^7 \Rightarrow m = 7$$

$$\text{And as virtual memory uses 6 bits} \Rightarrow n = 6$$

Address Translation

Page Table	
1	0
0	1
1	1
0	11
1	10
1	101
0	110
1	100
0	1000
1	11
0	1010
0	1011
0	1100
0	Null
0	Null
0	Null

How can we convert Virtual Address to Physical Address?

First let's see how VA and PA looks like:

- Virtual Address:

Virtual Page Number(VPN) (6-p bits)	Page Offset(PO) (p bits)
--	-----------------------------

- Physical Address:

Physical Page Number(PPN) (7-p bits)	Page Offset(PO) (p bits)
---	-----------------------------

How many bits do we need for VPN? That is how many bits do we require to index the page table?

Address Translation

Page Table	
1	0
0	1
1	1
0	11
1	10
1	101
0	110
1	100
0	1000
1	11
0	1010
0	1011
0	1100
0	Null
0	Null
0	Null

How can we convert Virtual Address to Physical Address?

First let's see how VA and PA looks like:

- Virtual Address:

Virtual Page Number(VPN) (6-p bits)	Page Offset(PO) (p bits)
--	-----------------------------

- Physical Address:

Physical Page Number(PPN) (7-p bits)	Page Offset(PO) (p bits)
---	-----------------------------

How many bits do we need for VPN? That is how many bits do we require to index the page table?

As page table has 16 PTE, so we need $16 = 2^4$ index, hence we will need 4 bits.

Address Translation

Page Table	
1	0
0	1
1	1
0	11
1	10
1	101
0	110
1	100
0	1000
1	11
0	1010
0	1011
0	1100
0	Null
0	Null
0	Null

How can we convert Virtual Address to Physical Address?

First let's see how VA and PA looks like:

- Virtual Address:

Virtual Page Number(VPN) (6-p bits)	Page Offset(PO) (p bits)
--	-----------------------------

- Physical Address:

Physical Page Number(PPN) (7-p bits)	Page Offset(PO) (p bits)
---	-----------------------------

How many bits do we need for VPN? That is how many bits do we require to index the page table?

As page table has 16 PTE, so we need $16 = 2^4$ index, hence we will need 4 bits.

So VPN will require 4 bits $\Rightarrow 6-p = 4 \Rightarrow p = 2$


Address Translation

Page Table	
1	0
0	1
1	1
0	11
1	10
1	101
0	110
1	100
0	1000
1	11
0	1010
0	1011
0	1100
0	Null
0	Null
0	Null


How can we convert Virtual Address to Physical Address?

First let's see how VA and PA looks like:

- Virtual Address:



- Physical Address:



How many bits do we need for VPN? That is how many bits do we require to index the page table?

As page table has 16 PTE, so we need $16 = 2^4$ index, hence we will need 4 bits.

So VPN will require 4 bits $\Rightarrow 6-p = 4 \Rightarrow p = 2$

Address Translation

Page Table	
1	0
0	1
1	1
0	11
1	10
1	101
0	110
1	100
0	1000
1	11
0	1010
0	1011
0	1100
0	Null
0	Null
0	Null

How can we convert Virtual Address to Physical Address?

First let's see how VA and PA looks like:

- Virtual Address:

Virtual Page Number(VPN) (4 bits)	Page Offset(PO) (2 bits)
--------------------------------------	-----------------------------

- Physical Address:

Physical Page Number(PPN) (5 bits)	Page Offset(PO) (2 bits)
---------------------------------------	-----------------------------

What will be 010010 in PA?

Page Table	
1	0
0	1
1	1
0	11
1	10
1	101
0	110
1	100
0	1000
1	11
0	1010
0	1011
0	1100
0	Null
0	Null
0	Null

So we have:

Virtual Page Number(VPN) (4 bits)	Page Offset(PO) (2 bits)
?	?

Physical Page Number(PPN) (5 bits)	Page Offset(PO) (2 bits)
?	?

What will be 010010 in PA?

Page Table	
1	0
0	1
1	1
0	11
1	10
1	101
0	110
1	100
0	1000
1	11
0	1010
0	1011
0	1100
0	Null
0	Null
0	Null

So we have:

Virtual Page Number(VPN) (4 bits)	Page Offset(PO) (2 bits)
0100	10

Physical Page Number(PPN) (5 bits)	Page Offset(PO) (2 bits)
?	?

What will be 010010 in PA?

Page Table	
1	0
0	1
1	1
0	11
1	10
1	101
0	110
1	100
0	1000
1	11
0	1010
0	1011
0	1100
0	Null
0	Null
0	Null

So we have:

Virtual Page Number(VPN) (4 bits)	Page Offset(PO) (2 bits)
0100	10

In decimal is 4

Physical Page Number(PPN) (5 bits)	Page Offset(PO) (2 bits)
?	?

What will be 010010 in PA?

Page Table	
1	0
0	1
1	1
0	11
1	10
1	101
0	110
1	100
0	1000
1	11
0	1010
0	1011
0	1100
0	Null
0	Null
0	Null

So we have:

Virtual Page Number(VPN) (4 bits)	Page Offset(PO) (2 bits)
0100	10

In decimal is 4

Physical Page Number(PPN) (5 bits)	Page Offset(PO) (2 bits)
?	?

What will be 010010 in PA?

Page Table	
1	0
0	1
1	1
0	11
1	10
1	101
0	110
1	100
0	1000
1	11
0	1010
0	1011
0	1100
0	Null
0	Null
0	Null

So we have:

Virtual Page Number(VPN) (4 bits)	Page Offset(PO) (2 bits)
0100	10

In decimal is 4

Physical Page Number(PPN) (5 bits)	Page Offset(PO) (2 bits)
?	?

Valid bit is 1, means it's cached

What will be 010010 in PA?

Page Table	
1	0
0	1
1	1
0	11
1	10
1	101
0	110
1	100
0	1000
1	11
0	1010
0	1011
0	1100
0	Null
0	Null
0	Null

So we have:

Virtual Page Number(VPN) (4 bits)	Page Offset(PO) (2 bits)
0100	10

In decimal is 4

Physical Page Number(PPN) (5 bits)	Page Offset(PO) (2 bits)
00010	?

What will be 010010 in PA?

Page Table	
1	0
0	1
1	1
0	11
1	10
1	101
0	110
1	100
0	1000
1	11
0	1010
0	1011
0	1100
0	Null
0	Null
0	Null

So we have:

Virtual Page Number(VPN) (4 bits)	Page Offset(PO) (2 bits)
0100	10

In decimal is 4

Physical Page Number(PPN) (5 bits)	Page Offset(PO) (2 bits)
00010	10

What will be 010010 in PA?

Page Table	
1	0
0	1
1	1
0	11
1	10
1	101
0	110
1	100
0	1000
1	11
0	1010
0	1011
0	1100
0	Null
0	Null
0	Null

So we have:

Virtual Page Number(VPN) (4 bits)	Page Offset(PO) (2 bits)
0100	10

In decimal is 4

Physical Page Number(PPN) (5 bits)	Page Offset(PO) (2 bits)
00010	10

So, PA will be: 0b0001010

What will be 000100 in PA?

Page Table	
1	0
0	1
1	1
0	11
1	10
1	101
0	110
1	100
0	1000
1	11
0	1010
0	1011
0	1100
0	Null
0	Null
0	Null

So we have:

Virtual Page Number(VPN) (4 bits)	Page Offset(PO) (2 bits)
?	?

Physical Page Number(PPN) (5 bits)	Page Offset(PO) (2 bits)
?	?

What will be 000110 in PA?

Page Table	
1	0
0	1
1	1
0	11
1	10
1	101
0	110
1	100
0	1000
1	11
0	1010
0	1011
0	1100
0	Null
0	Null
0	Null

So we have:

Virtual Page Number(VPN) (4 bits)	Page Offset(PO) (2 bits)
0001	00

Physical Page Number(PPN) (5 bits)	Page Offset(PO) (2 bits)
?	?

What will be 000110 in PA?

Page Table	
1	0
0	1
1	1
0	11
1	10
1	101
0	110
1	100
0	1000
1	11
0	1010
0	1011
0	1100
0	Null
0	Null
0	Null

So we have:

Virtual Page Number(VPN) (4 bits)	Page Offset(PO) (2 bits)
0001	00

In decimal is 1

Physical Page Number(PPN) (5 bits)	Page Offset(PO) (2 bits)
?	?

What will be 000110 in PA?

Page Table	
1	0
0	1
1	1
0	11
1	10
1	101
0	110
1	100
0	1000
1	11
0	1010
0	1011
0	1100
0	Null
0	Null
0	Null

So we have:

Virtual Page Number(VPN) (4 bits)	Page Offset(PO) (2 bits)
0001	00

In decimal is 1

Physical Page Number(PPN) (5 bits)	Page Offset(PO) (2 bits)
?	?

What will be 000110 in PA?

Page Table	
1	0
0	1
1	1
0	11
1	10
1	101
0	110
1	100
0	1000
1	11
0	1010
0	1011
0	1100
0	Null
0	Null
0	Null

So we have:

Virtual Page Number(VPN) (4 bits)	Page Offset(PO) (2 bits)
0001	00

In decimal is 1

Physical Page Number(PPN) (5 bits)	Page Offset(PO) (2 bits)
?	?

Valid bit is 0, means it's either Un-Cached, or Un-Allocated. In both cases we will get **Page Fault!**